

Interview Task — Real-time WebRTC VLM Multi-Object Detection (Phone → Browser → Inference → Overlay)

One-line goal: Build a reproducible demo that performs **real-time multi-object detection** on live video streamed from a phone via **WebRTC**, returns detection bounding boxes + labels to the browser, overlays them in near real-time, and deliver a **1-minute Loom video** showing the live demo, metrics, and one-sentence tradeoffs.

Deliverables (exact)

1. Git repo (frontend + optional server) with Dockerfile(s) and `docker-compose.yml` for local run. Include `start.sh` convenience script.
 2. `README.md` with one-command start instructions and mode-switch (server-mode / wasm-mode). Clear phone-join instructions (QR or short URL).
 3. `metrics.json` produced by a short bench run (30s) listing median & P95 end-to-end latency, processed FPS, and uplink/downlink kbps.
 4. A **1-minute Loom** video (hosted link) that: (a) shows phone → browser live overlay, (b) shows metrics output briefly, and (c) one-line improvement you'd do next.
 5. Short report (README appendix or `report.md`, 1 page) explaining design choices, low-resource mode, and backpressure policy.
-

Non-functional constraints & fairness

- **Low-resource path required.** Candidate must provide a mode runnable on modest laptops (no GPU). Typical approaches: WASM on-device inference (onnxruntime-web or tfjs-wasm), quantized small models, downscale input to 320×240, and adaptive sampling (10–15 FPS).
 - **Real-time from a phone.** Phone must use only a browser (Chrome on Android, Safari on iOS) to connect — no custom native app requirement.
 - **One command to start.** Candidate should supply `docker-compose up` or `./start.sh` to launch demo locally.
-

Minimal acceptance criteria (pass/fail)

- Phone can connect via QR/URL and stream live camera to the demo. Browser shows live overlays of bounding boxes aligned to frames.
 - `metrics.json` with median & P95 latency and FPS exists.
 - README explains how to run both low-resource and server modes.
 - Loom video demonstrates the live phone stream and metrics clearly within 1 minute.
-

UX / API contract (frame alignment)

Use this JSON message per frame over DataChannel / WebSocket for detection results (server → client):

```
{
  "frame_id": "string_or_int",
  "capture_ts": 1690000000000,
  "recv_ts": 1690000000100,
  "inference_ts": 1690000000120,
  "detections": [
    { "label": "person", "score": 0.93, "xmin": 0.12, "ymin": 0.08, "xmax": 0.34, "ymax": 0.67 }
  ]
}
```

- Coordinates normalized [0..1] to simplify overlay across resolutions.
 - Browser uses `capture_ts` and `frame_id` to align overlays with the correct frame and compute E2E latency.
-

Measurement & bench instructions

- **E2E latency (per frame):** `overlay_display_ts - capture_ts` → report median & P95 over a 30s run.
- **Server latency:** `inference_ts - recv_ts`.
- **Network latency:** `recv_ts - capture_ts`.
- **Processed FPS:** count of frames with detections displayed / seconds.
- **Bandwidth:** estimate via browser network inspector or tools like `ifstat/nethogs` during run.

Provide a simple bench script `./bench/run_bench.sh --duration 30 --mode server` that outputs `metrics.json`.

Low-resource guidance (what candidate *must* provide)

- **WASM on-device mode** using `onnxruntime-web` or `tfjs-wasm` with a small quantized model (example: MobileNet-SSD or YOLOv5n quantized).
- **Downscale**: default input size 320×240 and target processing 10–15 FPS.
- **Frame thinning**: process only latest frames; maintain a fixed-length queue and drop old frames when overloaded.
- **Simple mode switch**: `MODE=wasm` vs `MODE=server` in `start.sh`.

Candidates must document CPU usage on a modest laptop (e.g., Intel i5, 8GB RAM) for both modes.

Suggested technology & third-party components (install on dev machine OR phone)

For phone (user-facing, minimal required)

- **Chrome (Android)** — recommended: stable Chrome app.
- **Safari (iOS)** — iOS Safari supports WebRTC but feature parity varies; recommend latest iOS.
- No app installs required — phone uses browser to open a QR/URL and stream.

Optional phone tools (only if candidate documents and uses them):

- **ngrok** or **localtunnel** — for exposing localhost to the phone if Wi-Fi NAT blocks direct connect (candidate must include free-tier instructions).
- **Termux (Android)** — optional for advanced phone-side testing (not required).

For dev laptop / server (recommended installs)

- **Docker & Docker Compose** — recommended for reproducible local environment.
- **Node.js (>=16)** — for frontend dev server and lightweight WebRTC gateway if used.
- **Python 3.9+** — if using `aiortc` or server-side Python inference.
- **ONNX Runtime**: `onnxruntime` (CPU) for server-mode; `onnxruntime-web` for browser WASM.

- **tfjs** (optional) — `@tensorflow/tfjs` or `tfjs-backend-wasm` for JS inference.
- **aiortc** (Python) or **pion** (Go) or **mediasoup** (Node) — pick any for a gateway that can receive WebRTC tracks.
- **ngrok** (optional) — for quick phone connectivity.

Model & assets

- **ONNX Model Zoo** — MobileNet-SSD, YOLO variants, or quantized models.
- **TensorFlow Lite models** — if using TF.js or tflite-web.

Tools for measurement & debugging

- **Chrome DevTools (webrtc-internals)** — inspect RTP stats.
 - **getStats()** WebRTC API — for per-RTCPeerConnection metrics.
 - **ifstat / iftop / nethogs** — bandwidth during run.
 - **ps/top/htop** — CPU & memory.
 - **tc (linux)** — simulate packet loss/latency for robustness tests (optional).
-

Step-by-step candidate run instructions (to include in README)

1. `git clone <repo>`
2. `./start.sh` (defaults to `MODE=wasm` if no GPU) or `docker-compose up --build`
3. Open `http://localhost:3000` on your laptop; scan displayed QR with your phone.
4. Allow camera on phone; you should see phone video mirrored on the laptop with overlays.
5. Run `./bench/run_bench.sh --duration 30 --mode wasm` to collect metrics; inspect `metrics.json`.

If phone cannot reach laptop directly: run `./start.sh --ngrok` to start ngrok and copy the public URL to the phone.

Troubleshooting tips (include these in README)

- If phone won't connect: ensure phone and laptop are on same network OR use `ngrok/localtunnel`.
- If overlays are misaligned: confirm timestamps (`capture_ts`) are being echoed and units match (ms).

- If CPU is high: reduce resolution to 320×240 or offload to WASM mode.
 - Use Chrome [webrtc-internals](#) to inspect packet send/receive times and jitter.
-

Quick evaluation rubric (one-liner)

- **Functionality (30%)**: phone stream + overlays + metrics exist.
 - **Latency (25%)**: median & p95 E2E latency are sensible for chosen mode.
 - **Robustness (15%)**: queue/drop/backpressure strategy & low-resource mode.
 - **Docs & reproducibility (15%)**: clear README + [docker-compose](#) + 1-min Loom.
 - **Design reasoning (15%)**: tradeoffs and improvement plan.
-